

Hybrid Temporal Forecaster: An Advanced Regime-Aware Architecture for Energy Demand Prediction

1st Sahil Sarawgi
Roll No: 230153

2nd Ayush Gupta
Roll No: 230092

Abstract—Accurate temporal forecasting of energy demand is critical for modern power grid management, specifically when dealing with highly volatile load transitions. This report details the development, validation, and comprehensive comparative analysis of a Level-10 state-of-the-art Hybrid Temporal Forecaster. Moving beyond standard static mechanisms, the proposed architecture replaces external Gaussian Mixture Models (GMM) with an Internal Neural Regime Classifier and introduces a novel Multi-head Regime-Gated Attention (RGaA) mechanism. We benchmark this end-to-end model against established Support Vector Regression (SVR) and Time-Series Transformer baselines. Advanced regularization and training strategies, namely Time-Series MixUp and volatility-based Curriculum Learning, are deployed to enhance generalization. Furthermore, we conduct a rigorous mathematical failure analysis pinpointing exact edge-case breakdowns.

Index Terms—Temporal Forecasting, Regime-Gated Attention, Time-Series MixUp, Transformer, SVR, Curriculum Learning

I. INTRODUCTION

Time-series forecasting of power loads (measured in MW Load) poses a significant challenge due to volatile consumption patterns and complex temporal dependencies. Traditional predictive models and standard self-attention mechanisms often struggle to adapt seamlessly to abrupt regime shifts in the underlying data distribution. To address this, we propose an Advanced Hybrid Forecaster, a fully end-to-end joint learning architecture designed to dynamically adjust its attention based on current exogenous regime probabilities natively learned by the network. Through a rigorous comparative ablation framework, we evaluate the proposed model's performance, interpretability, and robustness against standard baselines.

II. DATASET AND FEATURE ENGINEERING

The models are trained and evaluated on an energy demand dataset tracking Real Energy Demand spanning from 2018 to 2021. To provide essential temporal context to the machine learning algorithms, comprehensive feature engineering was performed. The dataset consists of 9 specific features: continuous datetime parameters (hour, day of week, month, day of year), historical sequence lags (1-hour, 2-hour, and 24-hour lags), and critical rolling statistics (24-hour rolling mean and 24-hour rolling standard deviation).

For the attention-based architectures (Transformer and Hybrid), the continuous time-series data is segmented into se-

quences of length 24, enabling the models to process a standardized lookback window for each prediction step.

III. METHODOLOGY

A. Baseline Models

To evaluate the effectiveness of our proposed forecaster, we implemented two established comparative baselines:

- **Support Vector Regression (SVR):** A baseline sequence-to-target model leveraging a Radial Basis Function (RBF) kernel, scaled using a StandardScaler to map continuous temporal features into a higher-dimensional predictive space.
- **Time-Series Transformer:** A sequence-to-sequence model containing a Positional Encoding layer. Because attention mechanisms do not naturally process data sequentially, this layer injects sine and cosine-based positional information into the sequence before passing it through multi-head self-attention mechanisms and input projections.

B. Advanced Hybrid Forecaster

The core of our project is the Level-10 Advanced Hybrid Forecaster, which processes multi-scale convolutional feature extraction alongside a set of novel attention mechanisms. The key architectural improvements include:

- **Internal Neural Regime Classifier:** This component eliminates the need for an external Gaussian Mixture Model (GMM) by natively learning and classifying the underlying time-series regimes end-to-end.
- **Regime-Gated Attention (RGaA):** A novel architectural block that modulates standard self-attention weights based on exogenous regime probabilities. It passes the data through a sequential linear layer followed by a Sigmoid activation function to act as a probability gate.

C. Advanced Training Strategies

To ensure robust model convergence and prevent overfitting on volatile sequences, we implemented two Level-10 advanced training protocols:

- **Curriculum Learning Strategy:** The dataset is dynamically evaluated for sequence volatility using the engineered 24-hour rolling standard deviation feature. Train-

ing batches are then intelligently split into "Easy" (normal load) and "Hard" (extreme transition) curriculums.

- **Time-Series MixUp:** We applied an advanced linear interpolation regularization technique across sequences using a beta distribution parameterized with an alpha value of 0.2.

IV. EXPERIMENTAL RESULTS AND ANALYSIS

A. Visual Baseline Forecasting Comparisons

Baseline tracking was visualized over a 2-week forecast period against the actual real energy demand. The SVR Baseline utilizing the RBF Kernel successfully tracks the broader daily cyclical energy demand, though it exhibits rigidity during rapid micro-shifts.




Fig. 1. SVR Baseline (RBF Kernel) vs Actual Energy Demand over a 2-Week period. Shows the SVR predictions tracking the cyclical actual load.

Conversely, the Transformer Baseline utilizing Self-Attention captures high-frequency variances more effectively but is prone to localized overshooting when the underlying regime fundamentally changes without warning.




Fig. 2. Transformer (Self-Attention) vs Actual Energy Demand over a 2-Week period. Note the slight overshooting during volatile transitions.

B. Attention Map Interpretability

A key requirement of the Level-10 technical validation is interpretability. Analysis of the Regime-Aware Attention map confirms that the model's internally learned regime heads successfully modulate attention weights. Visualizing Head 0 over a 24-hour lookback window (T-24 to T-1) demonstrates the network dynamically applying distinct, high-density attention weights to highly relevant localized lags, proving the network learns rather than memorizes.




Fig. 3. Regime-Aware Attention Map (Head 0). The heatmap displays localized, high-density attention weights spanning the T-24 to T-1 lookback window.

C. Robustness to Sensor Noise

To validate the model's mathematical resilience to real-world edge conditions, we mapped a robustness curve injecting Gaussian sensor noise directly into the model inputs. At a baseline of 0.0 standard deviation noise, the Test Mean Absolute Error (MAE) is approximately 1028 MW. As noise severity artificially scales to a 0.30 standard deviation, the Test MAE scales proportionally to roughly 1098 MW. This confirms graceful degradation of the model's predictive capability rather than catastrophic failure.




Fig. 4. Model Robustness to Input Sensor Noise. The plot illustrates the Test MAE gracefully degrading from 1028 MW to 1098 MW as Gaussian Noise Standard Deviation scales to 0.30.

V. FAILURE ANALYSIS

A vital requirement of the Level-10 advanced architecture rubric is moving beyond basic inference metrics to locate where the model fundamentally breaks down mathematically. By conducting rigorous failure analysis, we actively searched

for anomalies and pinpointed a highly specific decoupling event on 2019-03-04 at 00:00:00. A zoomed-in assessment of this sudden spike visually isolates the exact datetime threshold where the mathematical constraints of the Hybrid forecaster completely decouple from the ground truth load, providing critical diagnostic feedback for future parameter tuning.




Fig. 5. Zoomed Failure Analysis: Sudden Spike at 2019-03-04 00:00:00. The vertical dashed line pinpoints the exact point of mathematical failure against the actual ground truth load.

VI. CONCLUSION

The proposed Hybrid Temporal Forecaster effectively addresses the shortcomings of static baseline models by gating temporal features through a neural regime classifier. Significantly enhanced by advanced regularization strategies like Time-Series Mixup and volatility-based Curriculum Learning, the architecture displays robust stability against artificial sensor noise. Critically, through deep failure analysis, we identified specific edge-case thresholds that break the model, completing the lifecycle of rigorous machine learning model development.

REFERENCES

- [1] Hybrid Temporal Forecaster Source Code: `advanced_hybrid_model.py`.
- [2] Training Protocols and Regularization: `train_advanced.py`.
- [3] Baseline implementations: `baseline_transformer.py`, `baseline_svr.py`.
- [4] Evaluation and ablation metrics: `compare_all_models.py`, `generate_final_plots.py`, `failure_analysis.py`.
- [5] Extracted Data Features: `featured_data.csv` and `featured_real_energy.csv`.